



L40S qualified Tensor backend — W128, L2, S257, R32 (synthetic resource study)


Training memory: none checkpoints




Training memory: loop checkpoints




Training latency: loop checkpoints



Physical cache



Decode peak allocation



Four-token decode graph latency

