




Decode memory: W512 / L4 / S4096 / R64



Full decoder with Tensor attention



Training memory: W512 / L2 / S1024 / R32



Shared-cache decode kernel: B1 / H8 / S4097

